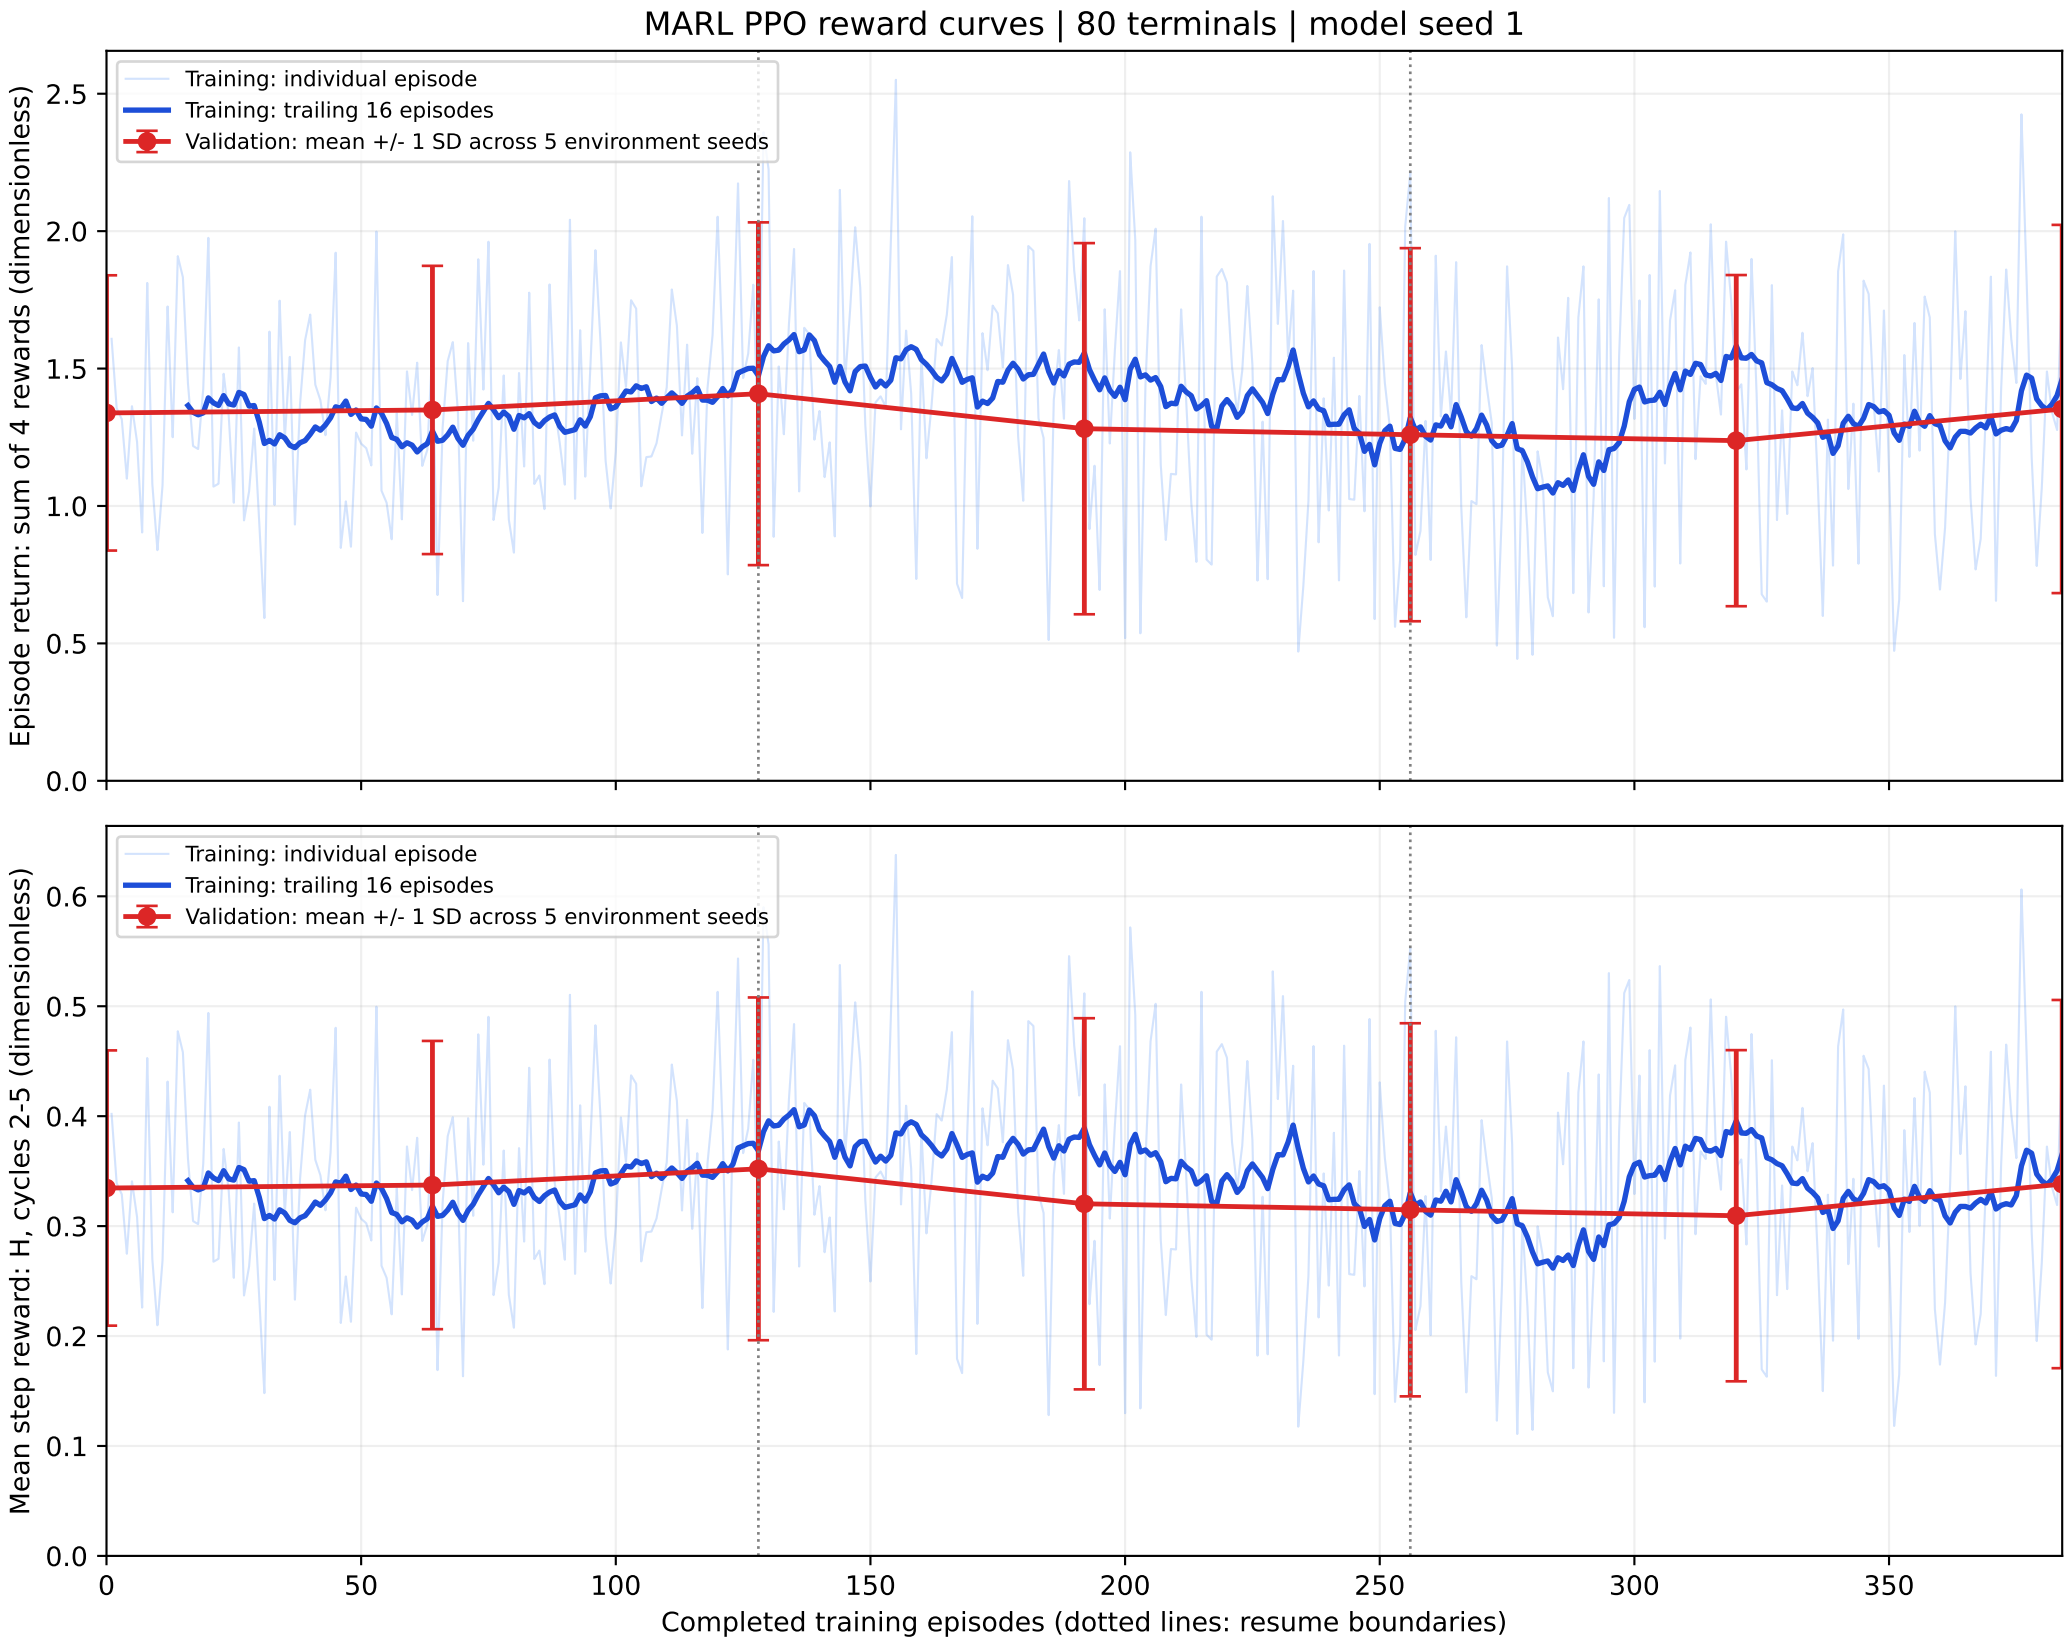
MARL PPO reward curves | 80 terminals | model seed 1
Training: individual episode
Training: trailing 16 episodes
Validation: mean +/- 1 SD across 5 environment seeds
Episode return: sum of 4 rewards (dimensionless)
Mean step reward: H, cycles 2-5 (dimensionless)
Completed training episodes (dotted lines: resume boundaries)
Reward = H_after; return = sum over cycles 2-5. Validation is unsmoothed. SD is not a confidence interval.