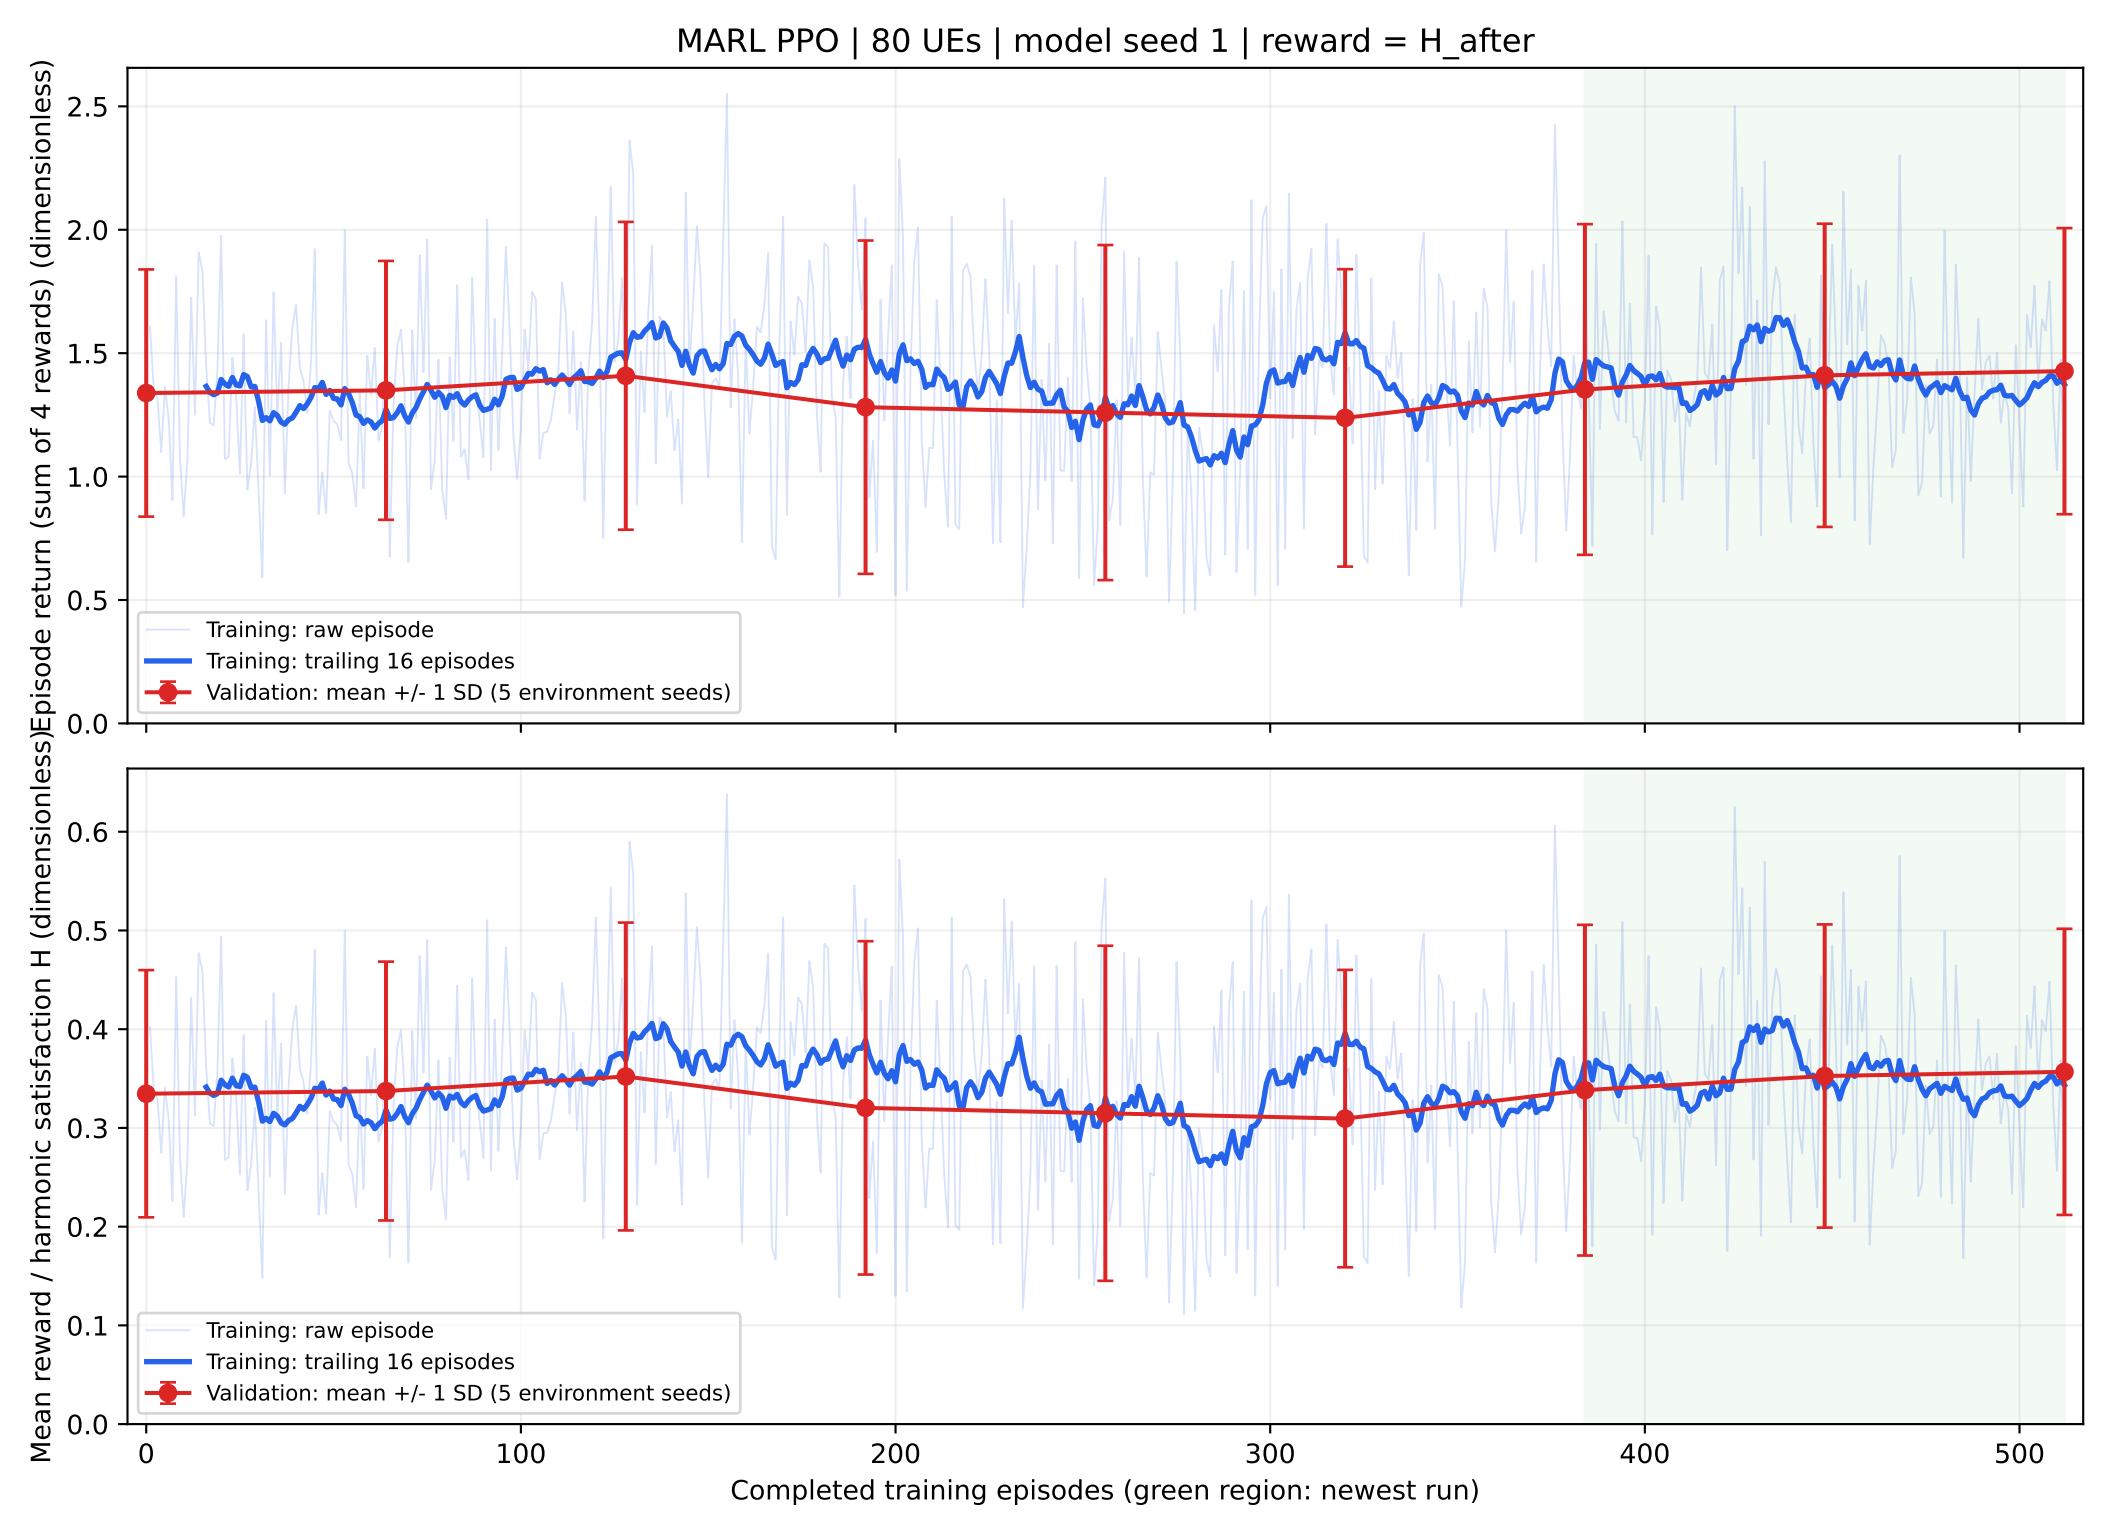
MARL PPO | 80 UEs | model seed 1 | reward = H_after
Episode return (sum of 4 rewards) (dimensionless)
Mean reward / harmonic satisfaction H (dimensionless)
Training: raw episode
Training: trailing 16 episodes
Validation: mean +/- 1 SD (5 environment seeds)
Completed training episodes (green region: newest run)