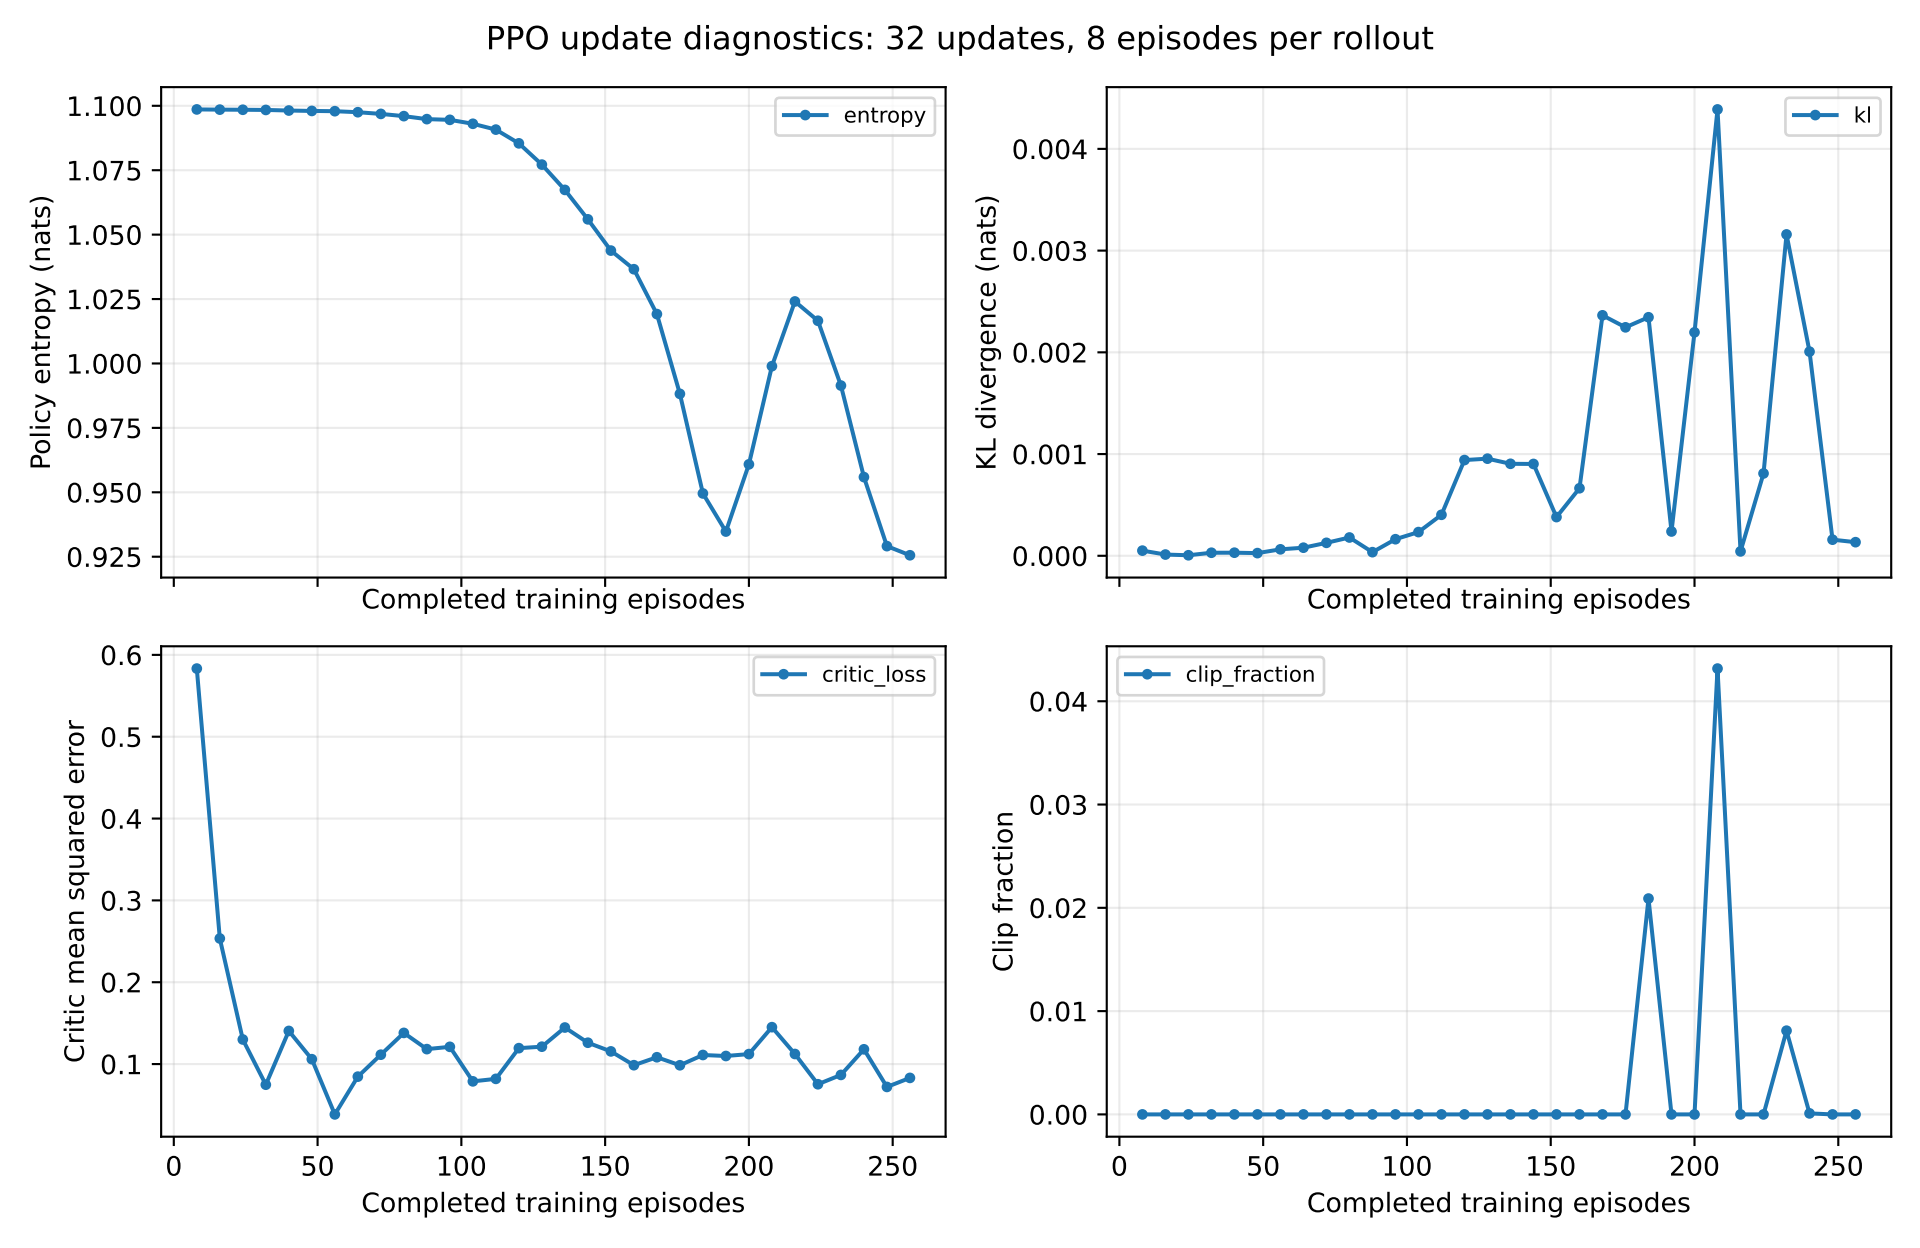
PPO update diagnostics: 32 updates, 8 episodes per rollout
entropy
Policy entropy (nats)
1.100
1.075
1.050
1.025
1.000
0.975
0.950
0.925
Completed training episodes
kl
KL divergence (nats)
0.004
0.003
0.002
0.001
0.000
Completed training episodes
critic_loss
Critic mean squared error
0.6
0.5
0.4
0.3
0.2
0.1
0
50
100
150
200
250
Completed training episodes
clip_fraction
Clip fraction
0.04
0.03
0.02
0.01
0.00
0
50
100
150
200
250
Completed training episodes